

Capstone 2: Stock Market Predictions

This capstone will complete all of the steps of the Data Science Method.

- 1) Problem identification
- 2) Data wrangling
- 3) Exploratory data analysis
- 4) Pre-processing and training data
- 5) Modelling
- 6) Documentation

1) Problem identification

I will be using the Kaggle dataset "S&P 500 stock data". The dataset is at <https://www.kaggle.com/datasets/camnugent/sandp500> and it covers 5 years of stocks including the date, open value, high value for that day, low value for that day, close value, volume, and stock name. You will see in the notebook but there are hundreds of stocks, and I will be cutting it to the first five stocks for proof of concept for the modeling process.

I will develop an accurate stock price prediction system in Python to predict the stock performance over a specific period. After doing feature engineering on the dataset, I can use ARMA, ARIMA and LSTM and see how those models do.

Problem statement formation

- Context: The stock market is extremely lucrative. Millions of people trade on it every day, and predicting movement has the potential for lots of profit and gain. It is volatile and hard to predict however, and it is a large industry of analysts on Wall Street to make good decisions for asset managers, banks, and other traders. Can time series analysis also be of use to predict future stock prices? Stock price might not be 100% dependent on past performance, of course, but we can make a model that is helpful.
- Criteria for success: This will be successful if it has more predictive power than a random walk and/or random noise.
- Scope of solution space: the solution space is stock trading advice on certain stocks, and will only depend on what we can infer from the time series prediction model, rather than other factors like what's going on with those companies and the market
- Constraints: we only have stock price information for a handful of stocks and for a certain period of time so we obviously can't predict outside of that
- Stakeholders: analysts, amateur traders,
- Data sources: kaggle dataset above, can reference <https://rpubs.com/kapage/523169> <https://www.projectpro.io/article/stock-price-prediction-using-machine-learning-project/571>

Stock Price as a Time Series Data

Treating stock data as time-series, one can use past stock prices (and other parameters) to predict the stock prices for the next day or week. Machine learning models such as

Recurrent Neural Networks (RNNs) or LSTMs are popular models applied to predicting time series data such as weather forecasting, election results, house prices, and, of course, stock prices. The idea is to weigh out the importance of recent and older data and determine which parameters affect the "current" or "next" day prices the most. The machine learning model assigns weights to each market feature and determines how much history the model should look at to predict future stock prices.

LSTM is a Recurrent Neural Network that works on data sequences, learning to retain only relevant information from a time window.

2) Data wrangling

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

sns.set_style('whitegrid')
plt.style.use("fivethirtyeight")
%matplotlib inline
```

```
In [2]: df=pd.read_csv("all_stocks_5yr.csv")
df.head()
```

```
Out[2]:
```

	date	open	high	low	close	volume	Name
0	2013-02-08	15.07	15.12	14.63	14.75	8407500	AAL
1	2013-02-11	14.89	15.01	14.26	14.46	8882000	AAL
2	2013-02-12	14.45	14.51	14.10	14.27	8126000	AAL
3	2013-02-13	14.30	14.94	14.25	14.66	10259500	AAL
4	2013-02-14	14.94	14.96	13.16	13.99	31879900	AAL

```
In [3]: earliest_date = df['date'].min()
latest_date = df['date'].max()

print("Earliest date: ", earliest_date)
print("Latest date: ", latest_date)
```

```
Earliest date: 2013-02-08
Latest date: 2018-02-07
```

Later on when I am breaking the data into training and test data, I will use the earliest ~80% of data as training.

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 619040 entries, 0 to 619039
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    date      619040 non-null  object
1    open      619029 non-null  float64
2    high      619032 non-null  float64
3    low       619032 non-null  float64
4    close     619040 non-null  float64
5    volume    619040 non-null  int64
6    Name      619040 non-null  object
dtypes: float64(4), int64(1), object(2)
memory usage: 33.1+ MB
```

```
In [5]: df.columns = df.columns.str.strip()
df.columns = df.columns.str.lower()
df.columns
```

```
Out[5]: Index(['date', 'open', 'high', 'low', 'close', 'volume', 'name'], dtype='object')
```

```
In [6]: # checking for missing values
print(df.isnull().sum())
```

```
date      0
open     11
high      8
low       8
close     0
volume    0
name      0
dtype: int64
```

Compared to the number of observations, 619,040, there are very few rows with missing values, at most 27 if the missing values are from all different rows. This is relatively negligible so I will drop these rows.

```
In [7]: df.dropna(inplace=True)
```

```
In [8]: print(df['name'].unique())
```

```
[ 'AAL' 'AAPL' 'AAP' 'ABBV' 'ABC' 'ABT' 'ACN' 'ADBE' 'ADI' 'ADM' 'ADP'
'ADSK' 'ADS' 'AEE' 'AEP' 'AES' 'AET' 'AFL' 'AGN' 'AIG' 'AIV' 'AIZ' 'AJG'
'AKAM' 'ALB' 'ALGN' 'ALK' 'ALLE' 'ALL' 'ALXN' 'AMAT' 'AMD' 'AME' 'AMGN'
'AMG' 'AMP' 'AMT' 'AMZN' 'ANDV' 'ANSS' 'ANTM' 'AON' 'AOS' 'APA' 'APC'
'APD' 'APH' 'APTV' 'ARE' 'ARNC' 'ATVI' 'AVB' 'AVGO' 'AVY' 'AWK' 'AXP'
'AYI' 'AZO' 'A' 'BAC' 'BAX' 'BA' 'BBT' 'BBY' 'BDX' 'BEN' 'BF.B' 'BHF'
'BHGE' 'BIIB' 'BK' 'BLK' 'BLL' 'BMY' 'BRK.B' 'BSX' 'BWA' 'BXP' 'CAG'
'CAH' 'CAT' 'CA' 'CBG' 'CBOE' 'CBS' 'CB' 'CCI' 'CCL' 'CDNS' 'CELG' 'CERN'
'CFG' 'CF' 'CHD' 'CHK' 'CHRW' 'CHTR' 'CINF' 'CI' 'CLX' 'CL' 'CMA' 'CMCSA'
'CME' 'CMG' 'CMI' 'CMS' 'CNC' 'CNP' 'COF' 'COG' 'COL' 'COO' 'COP' 'COST'
'COTY' 'CPB' 'CRM' 'CSCO' 'CSRA' 'CSX' 'CTAS' 'CTL' 'CTSH' 'CTXS' 'CVS'
'CVX' 'CXO' 'C' 'DAL' 'DE' 'DFS' 'DGX' 'DG' 'DHI' 'DHR' 'DISCA' 'DISCK'
'DISH' 'DIS' 'DLR' 'DLTR' 'DOV' 'DPS' 'DRE' 'DRI' 'DTE' 'DUK' 'DVA' 'DVN'
'DWDP' 'DXC' 'D' 'EA' 'EBAY' 'ECL' 'ED' 'EFX' 'EIX' 'EL' 'EMN' 'EMR'
'EOG' 'EQIX' 'EQR' 'EQT' 'ESRX' 'ESS' 'ES' 'ETFC' 'ETN' 'ETR' 'EVHC' 'EW'
'EXC' 'EXPD' 'EXPE' 'EXR' 'FAST' 'FBHS' 'FB' 'FCX' 'FDX' 'FE' 'FFIV'
'FISV' 'FIS' 'FITB' 'FLIR' 'FLR' 'FLS' 'FL' 'FMC' 'FOXA' 'FOX' 'FRT'
'FTI' 'FTV' 'F' 'GD' 'GE' 'GGP' 'GILD' 'GIS' 'GLW' 'GM' 'GOOGL' 'GOOG'
'GPC' 'GPN' 'GPS' 'GRMN' 'GS' 'GT' 'GWW' 'HAL' 'HAS' 'HBAN' 'HBI' 'HCA'
'HCN' 'HCP' 'HD' 'HES' 'HIG' 'HII' 'HLT' 'HOG' 'HOLX' 'HON' 'HPE' 'HPQ'
'HP' 'HRB' 'HRL' 'HRS' 'HSIC' 'HST' 'HSY' 'HUM' 'IBM' 'ICE' 'IDXX' 'IFF'
'ILMN' 'INCY' 'INFO' 'INTC' 'INTU' 'IPG' 'IP' 'IQV' 'IRM' 'IR' 'ISRG'
'ITW' 'IT' 'IVZ' 'JBHT' 'JCI' 'JEC' 'JNJ' 'JNPR' 'JPM' 'JWN' 'KEY' 'KHC'
'KIM' 'KLAC' 'KMB' 'KMI' 'KMX' 'KORS' 'KO' 'KR' 'KSS' 'KSU' 'K' 'LB'
'LEG' 'LEN' 'LH' 'LKQ' 'LLL' 'LLY' 'LMT' 'LNC' 'LNT' 'LOW' 'LRCX' 'LUK'
'LUV' 'LYB' 'L' 'MAA' 'MAC' 'MAR' 'MAS' 'MAT' 'MA' 'MCD' 'MCHP' 'MCK'
'MCO' 'MDLZ' 'MDT' 'MET' 'MGM' 'MHK' 'MKC' 'MLM' 'MMC' 'MMM' 'MNST' 'MON'
'MOS' 'MO' 'MPC' 'MRK' 'MRO' 'MSFT' 'MSI' 'MS' 'MTB' 'MTD' 'MU' 'MYL' 'M'
'NAVI' 'NBL' 'NCLH' 'NDAQ' 'NEE' 'NEM' 'NFLX' 'NFX' 'NI' 'NKE' 'NLSN'
'NOC' 'NOV' 'NRG' 'NSC' 'NTAP' 'NTRS' 'NUE' 'NVDA' 'NWL' 'NWSA' 'NWS'
'OKE' 'OMC' 'ORCL' 'ORLY' 'OXY' 'O' 'PAYX' 'PBCT' 'PCAR' 'PCG' 'PCLN'
'PDCO' 'PEG' 'PEP' 'PFE' 'PFG' 'PGR' 'PG' 'PHM' 'PH' 'PKG' 'PKI' 'PLD'
'PM' 'PNC' 'PNR' 'PNW' 'PPG' 'PPL' 'PRGO' 'PRU' 'PSA' 'PSX' 'PVH' 'PWR'
'PXD' 'PX' 'PYPL' 'QCOM' 'QRVO' 'RCL' 'REGN' 'REG' 'RE' 'RF' 'RHI' 'RHT'
'RJF' 'RL' 'RMD' 'ROK' 'ROP' 'ROST' 'RRC' 'RSG' 'RTN' 'SBAC' 'SBUX' 'SCG'
'SCHW' 'SEE' 'SHW' 'SIG' 'SJM' 'SLB' 'SLG' 'SNA' 'SNI' 'SNPS' 'SO' 'SPGI'
'SPG' 'SRCL' 'SRE' 'STI' 'STT' 'STX' 'STZ' 'SWKS' 'SWK' 'SYF' 'SYK'
'SYMC' 'SYI' 'TAP' 'TDG' 'TEL' 'TGT' 'TIF' 'TJX' 'TMK' 'TMO' 'TPR' 'TRIP'
'TROW' 'TRV' 'TSCO' 'TSN' 'TSS' 'TWX' 'TXN' 'TXT' 'T' 'UAA' 'UAL' 'UA'
'UDR' 'UHS' 'ULTA' 'UNH' 'UNM' 'UNP' 'UPS' 'URI' 'USB' 'UTX' 'VAR' 'VFC'
'VIAB' 'VLO' 'VMC' 'VNO' 'VRSK' 'VRSN' 'VRTX' 'VTR' 'VZ' 'V' 'WAT' 'WBA'
'WDC' 'WEC' 'WFC' 'WHR' 'WLTW' 'WMB' 'WMT' 'WM' 'WRK' 'WU' 'WYNN' 'WYN'
'WY' 'XEC' 'XEL' 'XLNX' 'XL' 'XOM' 'XRAY' 'XRX' 'XYL' 'YUM' 'ZBH' 'ZION'
'ZTS' ]
```

```
In [9]: len(df['name'].unique())
```

```
Out[9]: 505
```

These are all the stocks included.

```
In [10]: # checking for duplicate rows
df.duplicated().sum()
```

```
Out[10]: 0
```

Feature Engineering - making new features such as price change

```
In [11]: #https://www.kaggle.com/code/yassinesfaihi/preparing-the-data-for-stock-market-
#around line 59 there are more
df['price_change'] = df['close'] - df['open']
df['returns'] = df['close'].pct_change()
df['average_price'] = (df['close'] + df['open']) / 2
df['price_range'] = df['high'] - df['low']
```

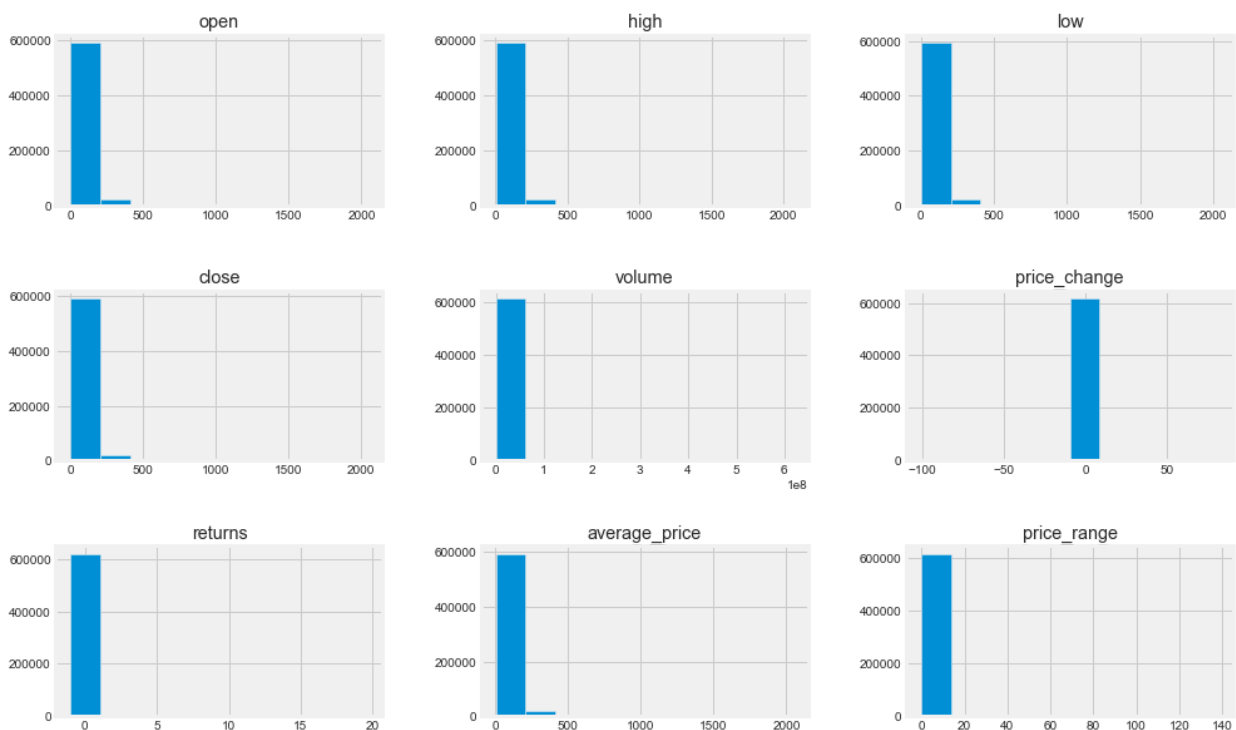
```
In [12]: df.describe()
```

```
Out[12]:
```

	open	high	low	close	volume	price_cl
count	619029.000000	619029.000000	619029.000000	619029.000000	6.190290e+05	619029.00
mean	83.023334	83.778419	82.256200	83.043305	4.321892e+06	0.0
std	97.378769	98.207735	96.507634	97.388913	8.693671e+06	1.55
min	1.620000	1.690000	1.500000	1.590000	1.010000e+02	-100.98
25%	40.220000	40.620000	39.830000	40.240800	1.070351e+06	-0.35
50%	62.590000	63.150000	62.020000	62.620000	2.082165e+06	0.02
75%	94.370000	95.180000	93.540000	94.410000	4.284550e+06	0.42
max	2044.000000	2067.990000	2035.110000	2049.000000	6.182376e+08	81.38

```
In [13]: #visualizing the data preliminarily
df.hist(figsize=(15,10))
plt.subplots_adjust(hspace=0.5);

#sns.pairplot(df)
#un-comment when it has time to load
```



Some are clustered towards zero, but maybe that's because of the scale/max being too high or an outlier. Let's visualize the outliers.

The IQR method is based on the interquartile range, or the difference between the 75th and 25th percentiles of the dataset. Data points that are outside of the range of 1.5IQR to 3IQR are considered outliers. I want to see the outliers of price range to consider getting rid of wildly fluctuating stocks in my training data.

```
In [14]: # Calculate the interquartile range (IQR) for the 'price' column
q1 = df['price_range'].quantile(0.25)
q3 = df['price_range'].quantile(0.75)
iqr = q3 - q1

# Calculate the lower and upper bounds for outliers
lower_bound = q1 - 1.5 * iqr
upper_bound = q3 + 1.5 * iqr

# Identify the outliers in the 'price' column
outliers = df[(df['price_range'] < lower_bound) | (df['price_range'] > upper_bound)]

# Print the outliers
print(outliers)
```

	date	open	high	low	close	volume	name	price_change
422	2014-10-13	31.07	31.3900	28.10	28.58	34532913	AAL	-2.49
425	2014-10-16	30.63	33.4000	30.00	32.97	24987103	AAL	2.34
430	2014-10-23	37.43	40.1800	36.80	38.48	33292004	AAL	1.05
467	2014-12-16	51.01	51.1500	47.68	47.96	22053666	AAL	-3.05
468	2014-12-17	48.26	49.4900	46.05	48.80	24779800	AAL	0.54
...
618420	2015-08-24	42.51	44.9700	37.73	43.53	5725519	ZTS	1.02
618461	2015-10-21	43.38	43.5600	39.53	41.66	9780195	ZTS	-1.72
618847	2017-05-04	55.48	59.6025	55.48	58.87	6923157	ZTS	3.39
618974	2017-11-02	64.30	67.8450	64.30	67.31	5185564	ZTS	3.01
619037	2018-02-05	76.64	76.9200	73.18	73.83	2962031	ZTS	-2.81


	returns	average_price	price_range
422	-0.071475	29.825	3.2900
425	0.040391	31.800	3.4000
430	0.038877	37.955	3.3800
467	-0.056648	49.485	3.4700
468	0.017515	48.530	3.4400
...
618420	-0.033096	43.020	7.2400
618461	-0.033186	42.520	4.0300
618847	0.059003	57.175	4.1225
618974	0.043404	65.805	3.5450
619037	-0.038421	75.235	3.7400

[49358 rows x 11 columns]

There are a lot of outliers because of the way I defined it with the IQR. I will choose to keep them in because it's a lot of valuable data.

```
In [15]: corr = df.corr()
fig, ax = plt.subplots(figsize=(10,10))
sns.heatmap(corr, linewidths=.5, ax=ax)
```

Out[15]: <AxesSubplot:>



The stock prices are correlated with themselves in a predictable way - the open and close for a single stock would be expected to be similar values.

3) Exploratory data analysis


We already did some EDA with pairplots and corr, but let's see some of the data more in depth.

```
In [16]: # set figure size
plt.figure(figsize=(11.7,8.27))

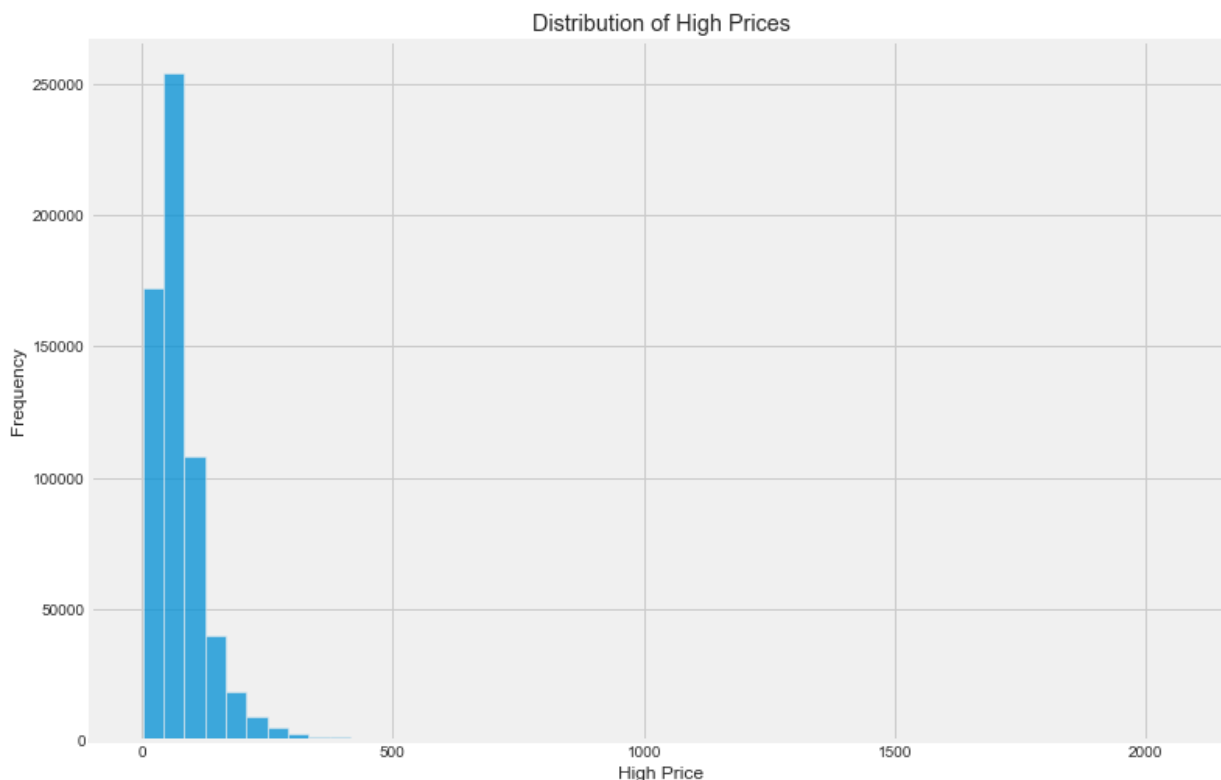
# create scatter plot
sns.scatterplot(x='open', y='close', data=df)

# add labels and title
plt.xlabel('Open Price')
plt.ylabel('Close Price')
plt.title('Open vs Close Price')
```

```
# show plot  
plt.show()
```



```
In [17]: # import necessary libraries  
import seaborn as sns  
import matplotlib.pyplot as plt  
  
# set figure size  
plt.figure(figsize=(11.7,8.27))  
  
# create histogram  
sns.histplot(data=df['high'], bins=50)  
  
# add labels and title  
plt.xlabel('High Price')  
plt.ylabel('Frequency')  
plt.title('Distribution of High Prices')  
  
# show plot  
plt.show()
```





```
In [18]: # set figure size
plt.figure(figsize=(11.7,8.27))

# create histogram
sns.histplot(data=df['low'], bins=50)

# add labels and title
plt.xlabel('Low Price')
plt.ylabel('Frequency')
plt.title('Distribution of Low Prices')

# show plot
plt.show()
```

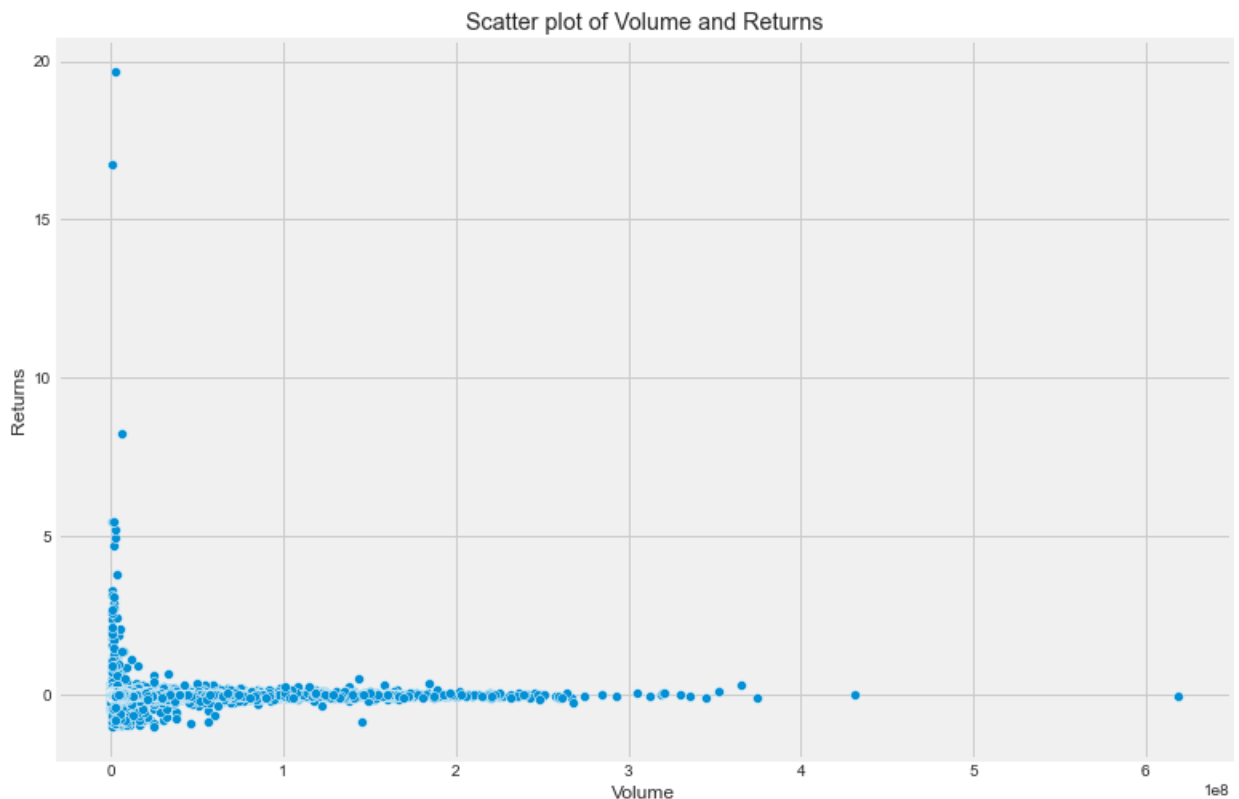


```
In [19]: # set figure size
plt.figure(figsize=(11.7,8.27))

# create scatter plot
sns.scatterplot(x='volume', y='returns', data=df)

# add labels and title
plt.xlabel('Volume')
plt.ylabel('Returns')
plt.title('Scatter plot of Volume and Returns')

# show plot
plt.show()
```



```
In [20]: #choosing 5 stocks to work with
# Create a list of the 5 stocks you want to include
selected_stocks = ['AAL', 'AAPL', 'AAP', 'ABBV', 'ABC']

# Filter the dataframe to only include rows where the 'Stock' column is in the
df = df.loc[df['name'].isin(selected_stocks)]

#there are 6295 rows now. everything works on the entire set too and if it was
```

Now it's time to train test split so we can work with the data to model.

```
In [21]: df.describe()
```

```
Out[21]:
```

	open	high	low	close	volume	price_change
count	6295.000000	6295.000000	6295.000000	6295.000000	6.295000e+03	6295.000000
mean	84.550957	85.396175	83.706077	84.570166	1.487854e+07	0.019209
std	40.196676	40.511310	39.861328	40.192771	2.518825e+07	1.214912
min	13.140000	13.420000	12.700000	13.020000	1.307120e+05	-9.110000
25%	52.760000	53.375000	52.220000	52.765000	1.584827e+06	-0.520000
50%	77.200000	77.840000	76.590000	77.150000	5.269007e+06	0.030000
75%	110.625000	111.875000	109.490000	110.410000	1.316181e+07	0.590000
max	201.240000	201.240000	198.160000	200.380000	2.668336e+08	13.890000

```
In [22]: df['date'] = pd.to_datetime(df['date']) #Convert the date column to datetime for
df_train = df[df['date'] < '2017-06-01']
```

```
df_test = df[df['date'] >= '2017-06-01']  
#change all df to df_train  
  
df_train = df_train.set_index('date') #set the date column as the index  
  
df_train = df_train.sort_index() #sort the dataframe by the index  
  
# Normalize the data (this is optional, but it can help improve the performance)  
from sklearn.preprocessing import MinMaxScaler  
  
scaler = MinMaxScaler()  
df_train[['open', 'high', 'low', 'close', 'volume', 'price_change', 'returns', 'ave
```

I didn't use train test split because it would have made the training random dates and not be helpful with predicting more random test set dates.

4) Pre-processing and training data

In [23]: `df_train.columns`

Out[23]: Index(['open', 'high', 'low', 'close', 'volume', 'name', 'price_change',
 'returns', 'average_price', 'price_range'],
 dtype='object')

In [24]: `df_test.columns`

Out[24]: Index(['date', 'open', 'high', 'low', 'close', 'volume', 'name',
 'price_change', 'returns', 'average_price', 'price_range'],
 dtype='object')

In [25]: `df`

Out [25]:

	date	open	high	low	close	volume	name	price_change	returns	average
0	2013-02-08	15.07	15.12	14.6300	14.75	8407500	AAL	-0.32	NaN	
1	2013-02-11	14.89	15.01	14.2600	14.46	8882000	AAL	-0.43	-0.019661	
2	2013-02-12	14.45	14.51	14.1000	14.27	8126000	AAL	-0.18	-0.013140	
3	2013-02-13	14.30	14.94	14.2500	14.66	10259500	AAL	0.36	0.027330	
4	2013-02-14	14.94	14.96	13.1600	13.99	31879900	AAL	-0.95	-0.045703	
...	
6290	2018-02-01	97.74	99.81	95.7300	99.29	2786798	ABC	1.55	-0.003813	
6291	2018-02-02	99.09	99.09	95.9100	96.02	1660267	ABC	-3.07	-0.032934	
6292	2018-02-05	95.62	96.52	91.6900	91.90	2278534	ABC	-3.72	-0.042908	
6293	2018-02-06	92.58	93.37	86.9403	91.54	4574997	ABC	-1.04	-0.003917	
6294	2018-02-07	91.60	95.34	91.1000	94.22	2509484	ABC	2.62	0.029277	

6295 rows × 11 columns

In [26]:

df_train

Out [26]:

	open	high	low	close	volume	name	price_change	returns	ave
date									
2013-02-08	0.010260	0.009051	0.010407	0.009234	0.031034	AAL	0.382174	NaN	
2013-02-08	0.290134	0.292735	0.292207	0.292668	0.592561	AAPL	0.402174	1.000000	
2013-02-08	0.177459	0.178229	0.182034	0.180775	0.004132	ABC	0.412174	0.083743	
2013-02-08	0.346624	0.352998	0.352151	0.351623	0.004377	AAP	0.420435	0.166395	
2013-02-08	0.123498	0.122458	0.124690	0.123986	0.051473	ABBV	0.390870	0.000000	
...	
2017-05-31	0.643275	0.649452	0.641216	0.643734	0.007453	AAP	0.373913	0.681624	
2017-05-31	0.748698	0.749388	0.753154	0.745837	0.091189	AAPL	0.343478	0.670786	
2017-05-31	0.186390	0.186668	0.187803	0.188888	0.016331	AAL	0.405217	0.686239	
2017-05-31	0.282297	0.282132	0.285291	0.282878	0.025516	ABBV	0.386522	0.676765	
2017-05-31	0.415683	0.417261	0.420845	0.420314	0.006522	ABC	0.415217	0.684973	

5425 rows × 10 columns

In [27]:

df_test

Out [27]:

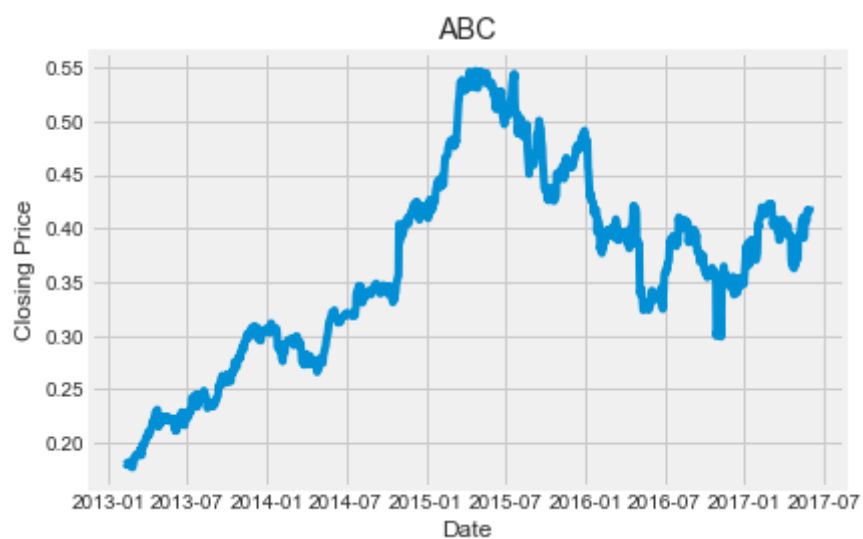


	date	open	high	low	close	volume	name	price_change	returns	average_
1085	2017-06-01	48.50	49.36	48.3000	49.05	4421404	AAL	0.55	0.013220	4
1086	2017-06-02	49.56	50.47	49.3700	49.52	7708567	AAL	-0.04	0.009582	4
1087	2017-06-05	49.53	49.95	49.4000	49.74	5466685	AAL	0.21	0.004443	4
1088	2017-06-06	49.47	50.10	49.3200	49.74	4473456	AAL	0.27	0.000000	4
1089	2017-06-07	49.44	50.91	48.8000	50.86	7078405	AAL	1.42	0.022517	!
...	
6290	2018-02-01	97.74	99.81	95.7300	99.29	2786798	ABC	1.55	-0.003813	9
6291	2018-02-02	99.09	99.09	95.9100	96.02	1660267	ABC	-3.07	-0.032934	9
6292	2018-02-05	95.62	96.52	91.6900	91.90	2278534	ABC	-3.72	-0.042908	9
6293	2018-02-06	92.58	93.37	86.9403	91.54	4574997	ABC	-1.04	-0.003917	9
6294	2018-02-07	91.60	95.34	91.1000	94.22	2509484	ABC	2.62	0.029277	9


870 rows × 11 columns

The X will be all the columns besides close, and close is the y. Here is each of the stocks in the training data plotted (should I be plotting for all time with df instead of df_train)?

```
In [28]: names = df_train['name'].unique() # get unique list of stock names
for name in names:
    # filter dataframe for each stock
    stock_df = df_train[df_train['name'] == name]

    # create plot
    plt.plot(stock_df.index, stock_df['close'])
    plt.title(name)
    plt.xlabel('Date')
    plt.ylabel('Closing Price')
    plt.show()
```





In these graphs the trend is mostly upward and non-stationary, but there is variation between the graphs. some seasonal fluctuations? there is also noise. additive or multiplicative?

5) Modelling

Model 1 - ARMA

```
In [29]: import statsmodels.api as sm
import itertools
from sklearn.metrics import mean_squared_error

import warnings
from statsmodels.tools.sm_exceptions import ConvergenceWarning
warnings.simplefilter('ignore', ConvergenceWarning)
warnings.filterwarnings("ignore")

# Create a dictionary to store the ARMA models for each stock
arma_models = {}

# Define the orders to test for the ARMA model
p = range(0, 3)
```

```

q = range(0, 3)
orders = list(itertools.product(p, q))

# Loop through each stock in the dataframe
for stock in df_train['name'].unique():

    # Filter the dataframe to only include the current stock
    stock_df = df_train[df_train['name'] == stock]['close']

    # Train multiple ARMA models with different orders
    best_aic = float('inf')
    best_order = None
    best_model = None

    for order in orders:
        try:
            arma_model = sm.tsa.SARIMAX(stock_df, order=(order[0], 0, order[1]))
            if arma_model.aic < best_aic:
                best_aic = arma_model.aic
                best_order = order
                best_model = arma_model
        except Exception as e:
            print(f"Error fitting ARMA model for {stock} with order {order}: {e}")
            continue

    # Store the best model in the dictionary with the stock name as the key
    arma_models[stock] = best_model

    # Print the summary of the best model
    if best_model is not None:
        print(f"Summary of best ARMA model for {stock}:")
        print(best_model.summary())

    # Predict the test data using the best model
    if stock in df_test['name'].unique():
        test_stock_df = df_test[df_test['name'] == stock]['close']
        pred = best_model.predict(start=len(df_train), end=len(df_train)+len(test_stock_df))
        mse_score = mean_squared_error(test_stock_df, pred)
        print(f"MSE score for {stock}: {mse_score}")

```

This problem is unconstrained.

Warning: more than 10 function and gradient evaluations in the last line search. Termination may possibly be caused by a bad search direction.
This problem is unconstrained.

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 1 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= -5.63177D-01 |proj g|= 5.26734D-04

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
1	1	15	1	0	0	5.267D-04	-5.632D-01

F = -0.56317684922351330

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= -5.63175D-01 |proj g|= 9.96082D-01

At iterate 5 f= -1.21033D+00 |proj g|= 3.25768D-01

At iterate 10 f= -1.23144D+00 |proj g|= 7.43214D-03

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
2	13	25	1	0	0	3.146D-08	-1.231D+00

F = -1.2314519411011695

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 3 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= -5.63171D-01 |proj g|= 9.93770D-01

At iterate 5 f= -1.73383D+00 |proj g|= 3.39349D+00

At iterate 10 f= -1.80066D+00 |proj g|= 9.12714D-03

This problem is unconstrained.

At iterate 15 f= -1.80085D+00 |proj g|= 1.45080D-01

At iterate 20 f= -1.80150D+00 |proj g|= 2.91669D-03

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
3	23	32	1	0	0	7.017D-07	-1.802D+00

F = -1.8015038721509726

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 4.45109D+02 |proj g|= 1.94835D+05

This problem is unconstrained.

```

At iterate    5    f= -2.21919D+00    |proj g|=  1.00354D+01
At iterate   10    f= -3.20703D+00    |proj g|=  9.18208D+00
At iterate   15    f= -3.89313D+00    |proj g|=  3.47092D+01
At iterate   20    f= -3.95637D+00    |proj g|=  1.29596D+00
At iterate   25    f= -3.96072D+00    |proj g|=  7.62511D-01
At iterate   30    f= -3.96089D+00    |proj g|=  2.24243D-05

```

* * *

```

Tit   = total number of iterations
Tnf   = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip  = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

* * *

```

      N      Tit      Tnf  Tnint  Skip  Nact      Projg      F
      2       32       71      1      0      0    2.776D-03  -3.961D+00
F = -3.9608871899371181

```

```

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
RUNNING THE L-BFGS-B CODE

```

* * *

```

Machine precision = 2.220D-16

```

```

N =          3      M =          10

```

```

At X0          0 variables are exactly at the bounds

```

```

At iterate    0    f=  3.96620D+02    |proj g|=  1.73892D+05

```

```

At iterate    5    f= -2.10116D+00    |proj g|=  6.54147D+00

```

This problem is unconstrained.

```

At iterate   10    f= -3.01409D+00    |proj g|=  8.21307D+00
At iterate   15    f= -3.01719D+00    |proj g|=  5.28515D-01
At iterate   20    f= -3.69171D+00    |proj g|=  3.10711D+01
At iterate   25    f= -3.93886D+00    |proj g|=  1.25370D+01
At iterate   30    f= -3.96008D+00    |proj g|=  4.24147D-01
At iterate   35    f= -3.96266D+00    |proj g|=  3.04308D-01

```

* * *

```

Tit   = total number of iterations
Tnf   = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip  = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

* * *

```

N      Tit      Tnf  Tnint  Skip  Nact      Projg      F
  3      39      96      1      0      0  2.354D-02  -3.963D+00
F = -3.9626799650870326

```

```

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
RUNNING THE L-BFGS-B CODE

```

* * *

```
Machine precision = 2.220D-16
```

```
N =          4      M =          10
```

```
At X0          0 variables are exactly at the bounds
```

```
At iterate    0    f=  3.68898D+02    |proj g|=  1.61850D+05
```

Warning: more than 10 function and gradient evaluations in the last line search. Termination may possibly be caused by a bad search direction. This problem is unconstrained.

```

At iterate    5    f= -2.01905D+00    |proj g|=  5.46834D+00
At iterate   10    f= -2.88812D+00    |proj g|=  5.47683D+00
At iterate   15    f= -3.01584D+00    |proj g|=  2.35858D+01
At iterate   20    f= -3.75926D+00    |proj g|=  4.27289D+00
At iterate   25    f= -3.76487D+00    |proj g|=  2.53707D+00
At iterate   30    f= -3.80299D+00    |proj g|=  2.40521D+01
At iterate   35    f= -3.92165D+00    |proj g|=  2.04802D+01
At iterate   40    f= -3.93757D+00    |proj g|=  1.88793D+01
At iterate   45    f= -3.96200D+00    |proj g|=  2.72766D+00
At iterate   50    f= -3.96338D+00    |proj g|=  3.83997D-01

```

* * *

```

Tit   = total number of iterations
Tnf   = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip  = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

* * *

```

      N      Tit      Tnf  Tnint  Skip  Nact      Projg      F
      4       50      85      1      0      0    3.840D-01  -3.963D+00
F = -3.9633840949170689

```

```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT
RUNNING THE L-BFGS-B CODE

```

* * *

```
Machine precision = 2.220D-16
```

```
N =          3      M =          10
```

```
At X0          0 variables are exactly at the bounds
```

```
At iterate    0    f=  4.46438D+02    |proj g|=  1.95700D+05
```

This problem is unconstrained.

```

At iterate    5    f= -2.32297D+00    |proj g|=  9.53873D+00
At iterate   10    f= -3.47017D+00    |proj g|=  2.57549D+01
At iterate   15    f= -3.61146D+00    |proj g|=  7.79588D-01
At iterate   20    f= -3.61217D+00    |proj g|=  1.46978D-02
At iterate   25    f= -3.61517D+00    |proj g|=  3.63556D+00
At iterate   30    f= -3.95417D+00    |proj g|=  7.02066D-01
At iterate   35    f= -3.96232D+00    |proj g|=  3.49144D-01
At iterate   40    f= -3.96280D+00    |proj g|=  1.27008D-01
At iterate   45    f= -3.96281D+00    |proj g|=  6.41147D-05

```

* * *

```

Tit   = total number of iterations
Tnf   = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip  = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

* * *

```

      N      Tit      Tnf  Tnint  Skip  Nact      Projg      F
      3      45      84      1      0      0      6.411D-05  -3.963D+00
F = -3.9628080378647197

```

```

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
RUNNING THE L-BFGS-B CODE

```

* * *

```

Machine precision = 2.220D-16
N = 3      M = 10

```

```

At X0      0 variables are exactly at the bounds

```

```

At iterate    0    f=  1.94402D+03    |proj g|=  8.46153D+05

```

This problem is unconstrained.


```

ys=-5.286E-01  -gs= 1.113E+00 BFGS update SKIPPED

At iterate    5    f= -3.88486D-01    |proj g|=  2.34549D+00
ys=-7.792E+00  -gs= 1.337E+00 BFGS update SKIPPED

At iterate   10    f= -3.65523D+00    |proj g|=  1.91960D+00

At iterate   15    f= -3.94174D+00    |proj g|=  4.35151D+00

At iterate   20    f= -3.96018D+00    |proj g|=  2.62834D-01

At iterate   25    f= -3.96078D+00    |proj g|=  1.51888D-01

At iterate   30    f= -3.96078D+00    |proj g|=  4.73424D-03

At iterate   35    f= -3.96078D+00    |proj g|=  7.62542D-02

At iterate   40    f= -3.96079D+00    |proj g|=  7.82758D-01

At iterate   45    f= -3.96089D+00    |proj g|=  1.66703D-01

    * * *

Tit   = total number of iterations
Tnf   = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip  = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

```

    * * *

      N    Tit    Tnf  Tnint  Skip  Nact    Projg    F
      4     48    104     1     2     0    1.233D-03  -3.961D+00
F = -3.9608892824787967

```

```

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
RUNNING THE L-BFGS-B CODE

```

```

    * * *

Machine precision = 2.220D-16
N =          5      M =          10

At X0          0 variables are exactly at the bounds

At iterate    0    f=  2.08335D+02    |proj g|=  9.21649D+04

```

Warning: more than 10 function and gradient evaluations in the last line search. Termination may possibly be caused by a bad search direction. This problem is unconstrained.

ys=-1.617E+00 -gs= 8.840E-01 BFGS update SKIPPED

```

At iterate    5    f= -1.03966D+00    |proj g|=  4.11516D+00
At iterate   10    f= -3.22850D+00    |proj g|=  1.52645D+00
At iterate   15    f= -3.53216D+00    |proj g|=  1.99493D+01
At iterate   20    f= -3.53663D+00    |proj g|=  1.36531D+00
At iterate   25    f= -3.54562D+00    |proj g|=  1.58203D+01
At iterate   30    f= -3.94877D+00    |proj g|=  1.43810D+01
At iterate   35    f= -3.95078D+00    |proj g|=  1.43779D+00
At iterate   40    f= -3.95465D+00    |proj g|=  2.52533D+00
At iterate   45    f= -3.95840D+00    |proj g|=  6.06711D-01
At iterate   50    f= -3.95849D+00    |proj g|=  2.60106D-03

```

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
5	50	91	1	1	0	2.601D-03	-3.958D+00

F = -3.9584937690586632

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT

Summary of best ARMA model for AAL:

SARIMAX Results

```

=====
Dep. Variable:          close    No. Observations:          1085
Model:                SARIMAX(2, 0, 0)    Log Likelihood          4299.647
Date:                Sun, 26 Mar 2023    AIC          -8593.293
Time:                15:52:43    BIC          -8578.325
Sample:                0    HQIC          -8587.627
                    - 1085

```

Covariance Type: opg

```

=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          1.0613      0.027     39.869      0.000        1.009        1.113
ar.L2         -0.0619      0.027     -2.325      0.020       -0.114       -0.010
sigma2        2.098e-05    6.22e-07    33.721      0.000       1.98e-05       2.22e-05
=====

```

```

=====
Ljung-Box (L1) (Q):                0.01    Jarque-Bera (JB):                2
52.97
Prob(Q):                0.91    Prob(JB):

```

```

0.00
Heteroskedasticity (H):          2.04   Skew:
-0.26
Prob(H) (two-sided):          0.00   Kurtosis:
5.31
=====
=====

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

MSE score for AAL: 2500.3323356590035

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 1 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 6.95168D-01 |proj g|= 4.25250D-05

This problem is unconstrained.

Warning: more than 10 function and gradient
evaluations in the last line search. Termination
may possibly be caused by a bad search direction.
This problem is unconstrained.

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
1	1	16	1	0	0	4.251D-05	6.952D-01

F = 0.69516781542909423

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 6.95169D-01 |proj g|= 9.96564D-01

At iterate 5 f= 7.81857D-02 |proj g|= 2.81223D-01

At iterate 10 f= 1.73546D-02 |proj g|= 6.21207D-02

At iterate 15 f= 1.40729D-02 |proj g|= 5.31804D-03

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
2	17	26	1	0	0	1.577D-05	1.407D-02

F = 1.4072072105632067E-002

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
 RUNNING THE L-BFGS-B CODE

* * *

This problem is unconstrained.

Machine precision = 2.220D-16

N = 3 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 6.95172D-01 |proj g|= 9.94408D-01

At iterate 5 f= -5.39876D-01 |proj g|= 4.58850D-01

At iterate 10 f= -5.90641D-01 |proj g|= 2.16742D-03

At iterate 15 f= -5.90676D-01 |proj g|= 4.16572D-02

Bad direction in the line search;
refresh the lbfgs memory and restart the iteration.

Warning: more than 10 function and gradient
evaluations in the last line search. Termination
may possibly be caused by a bad search direction.
This problem is unconstrained.

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
3	19	63	2	0	0	2.681D-04	-5.907D-01

F = -0.59068791256200581

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.89770D+03 |proj g|= 4.82636D+05

At iterate 5 f= 9.61300D-04 |proj g|= 1.26399D+00
 ys=-2.566E-01 -gs= 7.977E-01 BFGS update SKIPPED

At iterate 10 f= -3.40228D+00 |proj g|= 3.75506D+00

At iterate 15 f= -3.41186D+00 |proj g|= 7.95623D+00

At iterate 20 f= -3.42078D+00 |proj g|= 1.56461D+00

At iterate 25 f= -3.42152D+00 |proj g|= 1.87582D-01

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
2	29	65	1	1	0	2.743D-04	-3.422D+00

F = -3.4215358296276110

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 3 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.82465D+03 |proj g|= 4.63887D+05

Warning: more than 10 function and gradient
evaluations in the last line search. Termination
may possibly be caused by a bad search direction.
This problem is unconstrained.

At iterate 5 f= -1.25719D+00 |proj g|= 3.10693D+00

At iterate 10 f= -2.47564D+00 |proj g|= 1.23983D+01

At iterate 15 f= -2.60557D+00 |proj g|= 2.25598D+00

At iterate 20 f= -3.33351D+00 |proj g|= 6.52266D+01

At iterate 25 f= -3.41358D+00 |proj g|= 5.79943D-01

At iterate 30 f= -3.42050D+00 |proj g|= 3.72798D-01

At iterate 35 f= -3.42170D+00 |proj g|= 2.66385D-02

At iterate 40 f= -3.42178D+00 |proj g|= 5.70236D-04

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
3	40	78	1	0	0	5.702D-04	-3.422D+00

F = -3.4217782475748235

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 4 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.87401D+03 |proj g|= 4.75944D+05

This problem is unconstrained.

```
At iterate    5    f= -6.78481D-01    |proj g|=  4.91206D+00
At iterate   10    f= -2.18645D+00    |proj g|=  5.32706D+00
At iterate   15    f= -2.83060D+00    |proj g|=  2.82817D+00
At iterate   20    f= -3.32333D+00    |proj g|=  4.82425D+00
At iterate   25    f= -3.40366D+00    |proj g|=  2.52109D+00
At iterate   30    f= -3.41227D+00    |proj g|=  2.64470D-01
At iterate   35    f= -3.41737D+00    |proj g|=  8.88475D+00
At iterate   40    f= -3.42156D+00    |proj g|=  1.64763D+00
At iterate   45    f= -3.42183D+00    |proj g|=  2.25526D-01
```

Warning: more than 10 function and gradient
evaluations in the last line search. Termination
may possibly be caused by a bad search direction.
This problem is unconstrained.

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Proyg	F
4	48	100	1	0	0	2.450D-02	-3.422D+00

F = -3.4218294294125817

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 3 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.89696D+03 |proj g|= 4.82354D+05

At iterate 5 f= -1.03181D+00 |proj g|= 3.31669D+00

At iterate 10 f= -2.65325D+00 |proj g|= 2.58064D+01

At iterate 15 f= -3.06888D+00 |proj g|= 3.68620D+00

At iterate 20 f= -3.07548D+00 |proj g|= 1.15273D-02

At iterate 25 f= -3.07555D+00 |proj g|= 3.77603D-01

At iterate 30 f= -3.35754D+00 |proj g|= 3.67246D+01

At iterate 35 f= -3.41826D+00 |proj g|= 2.30300D+00

At iterate 40 f= -3.42163D+00 |proj g|= 4.15878D-02

At iterate 45 f= -3.42177D+00 |proj g|= 7.65864D-03

Warning: more than 10 function and gradient
 evaluations in the last line search. Termination
 may possibly be caused by a bad search direction.
 This problem is unconstrained.

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
3	46	104	1	0	0	7.659D-03	-3.422D+00

F = -3.4217731597268064

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 4 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 6.04586D+02 |proj g|= 1.54302D+05

At iterate 5 f= -1.50301D+00 |proj g|= 1.17936D+00

At iterate 10 f= -3.13580D+00 |proj g|= 2.58615D+01

At iterate 15 f= -3.36572D+00 |proj g|= 1.75047D+00

At iterate 20 f= -3.37806D+00 |proj g|= 1.30514D-01

At iterate 25 f= -3.39449D+00 |proj g|= 1.14040D+00

At iterate 30 f= -3.41439D+00 |proj g|= 1.64324D+00
 ys=-2.833E-01 -gs= 1.790E-04 BFGS update SKIPPED

At iterate 35 f= -3.41460D+00 |proj g|= 3.76100D-02

Bad direction in the line search;
 refresh the lbfgs memory and restart the iteration.

Line search cannot locate an adequate point after MAXLS
 function and gradient evaluations.

Previous x, f and g restored.

Possible causes: 1 error in function or gradient evaluation;
 2 rounding error dominate computation.

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
4	36	142	2	1	0	3.761D-02	-3.415D+00

F = -3.4145957632653183

ABNORMAL_TERMINATION_IN_LNSRCH
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 5 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 3.67485D+06 |proj g|= 9.31810D+08

This problem is unconstrained.

```

At iterate    5    f=  1.69197D+01    |proj g|=  3.96336D+01

At iterate   10    f=  1.93757D+00    |proj g|=  1.34899D+00

At iterate   15    f= -7.38877D-01    |proj g|=  8.86517D-01

At iterate   20    f= -1.26155D+00    |proj g|=  5.59920D+00
ys=-1.364E+00  -gs=  8.036E-01 BFGS update SKIPPED

At iterate   25    f= -3.34339D+00    |proj g|=  9.39311D+00

At iterate   30    f= -3.37319D+00    |proj g|=  5.10216D+00

At iterate   35    f= -3.41420D+00    |proj g|=  1.62455D-01

At iterate   40    f= -3.41757D+00    |proj g|=  5.75209D-01

At iterate   45    f= -3.41769D+00    |proj g|=  3.06491D-01

```

* * *

```

Tit   = total number of iterations
Tnf   = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip  = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

* * *

```

      N      Tit      Tnf  Tnint  Skip  Nact      Projg      F
      5       47       93       1      1      0    1.186D-01  -3.418D+00
F = -3.4176971114059382

```

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH

Summary of best ARMA model for AAPL:

SARIMAX Results

```

=====
Dep. Variable:          close    No. Observations:          1085
Model:                SARIMAX(1, 0, 0)    Log Likelihood          3712.366
Date:                Sun, 26 Mar 2023    AIC                    -7420.733
Time:                15:52:50    BIC                    -7410.754
Sample:                0    HQIC                    -7416.955
                    - 1085

```

Covariance Type: opg

```

=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.9999          0.000    4235.288      0.000          0.999          1.000
sigma2         6.189e-05    1.59e-06     38.812      0.000      5.88e-05      6.5e-05
=====

```

```

=====
Ljung-Box (L1) (Q):                0.40    Jarque-Bera (JB):                5
86.54
Prob(Q):                0.53    Prob(JB):
0.00
Heteroskedasticity (H):            1.85    Skew:
-0.12
Prob(H) (two-sided):            0.00    Kurtosis:

```

6.59

```
=====
=====
```

Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients (complex
-step).
```

```
MSE score for AAPL: 25897.155535940266
```

```
RUNNING THE L-BFGS-B CODE
```

```
* * *
```

```
Machine precision = 2.220D-16
```

```
N =          1      M =          10
```

```
At X0          0 variables are exactly at the bounds
```

```
At iterate    0      f=  4.39471D-01      |proj g|=  7.09145D-05
```

```
* * *
```

```
Tit   = total number of iterations
```

```
Tnf   = total number of function evaluations
```

```
Tnint = total number of segments explored during Cauchy searches
```

```
Skip  = number of BFGS updates skipped
```

```
Nact  = number of active bounds at final generalized Cauchy point
```

```
Projg = norm of the final projected gradient
```

```
F     = final function value
```

```
* * *
```

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
1	1	17	1	0	0	7.091D-05	4.395D-01

F = 0.43947114509010526

```
CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
```

```
RUNNING THE L-BFGS-B CODE
```

```
* * *
```

```
Machine precision = 2.220D-16
```

```
N =          2      M =          10
```

```
At X0          0 variables are exactly at the bounds
```

```
At iterate    0      f=  7.34704D-02      |proj g|=  1.95942D+00
```

```
At iterate    5      f= -2.26670D-01      |proj g|=  2.24359D-01
```

```
At iterate   10      f= -2.41771D-01      |proj g|=  5.63611D-02
```

This problem is unconstrained.

Warning: more than 10 function and gradient evaluations in the last line search. Termination may possibly be caused by a bad search direction.
This problem is unconstrained.

At iterate 15 f= -2.42608D-01 |proj g|= 1.14770D-02

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Proyg	F
2	19	27	1	0	0	7.349D-06	-2.426D-01

F = -0.24262004509623164

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16
 N = 3 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 4.74293D-03 |proj g|= 2.34037D+00

At iterate 5 f= -6.40629D-01 |proj g|= 2.25224D-01

This problem is unconstrained.

At iterate 10 f= -6.82793D-01 |proj g|= 2.72141D+00

At iterate 15 f= -7.56126D-01 |proj g|= 6.67397D-04

At iterate 20 f= -7.56128D-01 |proj g|= 2.50484D-02

At iterate 25 f= -7.56322D-01 |proj g|= 1.64471D-01

At iterate 30 f= -7.56514D-01 |proj g|= 7.45930D-04

Bad direction in the line search;
 refresh the lbfgs memory and restart the iteration.

Line search cannot locate an adequate point after MAXLS
 function and gradient evaluations.

Previous x, f and g restored.

Possible causes: 1 error in function or gradient evaluation;
 2 rounding error dominate computation.

This problem is unconstrained.

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
3	32	89	2	0	0	1.146D-04	-7.565D-01

F = -0.75651389655353907

ABNORMAL_TERMINATION_IN_LNSRCH
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.06724D+03 |proj g|= 7.08137D+05

At iterate 5 f= -2.58322D+00 |proj g|= 2.38301D+01

At iterate 10 f= -3.47577D+00 |proj g|= 1.15388D+02

At iterate 15 f= -3.68612D+00 |proj g|= 3.77684D+00

At iterate 20 f= -3.71852D+00 |proj g|= 4.57479D-01

At iterate 25 f= -3.72099D+00 |proj g|= 4.62438D-01

At iterate 30 f= -3.72118D+00 |proj g|= 4.78361D-04

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
2	31	72	1	0	0	2.397D-05	-3.721D+00

F = -3.7211789417505607

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 3 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.09754D+03 |proj g|= 7.17576D+05

This problem is unconstrained.

At iterate 5 f= -1.28467D+00 |proj g|= 1.16640D+00

At iterate 10 f= -2.70103D+00 |proj g|= 5.12051D+00

At iterate 15 f= -3.16217D+00 |proj g|= 3.41660D+00

At iterate 20 f= -3.18779D+00 |proj g|= 8.57610D+00

At iterate 25 f= -3.70973D+00 |proj g|= 2.12382D+01

At iterate 30 f= -3.71944D+00 |proj g|= 3.76619D-01

At iterate 35 f= -3.72056D+00 |proj g|= 3.73904D+00

At iterate 40 f= -3.72120D+00 |proj g|= 7.87553D-01

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
3	44	97	1	0	0	8.070D-03	-3.721D+00

F = -3.7212082416370458

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 4 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.96666D+03 |proj g|= 6.72170D+05

Warning: more than 10 function and gradient
evaluations in the last line search. Termination
may possibly be caused by a bad search direction.

This problem is unconstrained.


```

At iterate    5    f= -1.49474D+00    |proj g|=  5.28693D+00
At iterate   10    f= -2.88671D+00    |proj g|=  2.51914D+00
At iterate   15    f= -2.96518D+00    |proj g|=  4.22387D-01
At iterate   20    f= -3.26474D+00    |proj g|=  2.20388D+01
At iterate   25    f= -3.61798D+00    |proj g|=  8.64903D-01
At iterate   30    f= -3.63013D+00    |proj g|=  1.31105D+00
At iterate   35    f= -3.64888D+00    |proj g|=  8.89536D+00
At iterate   40    f= -3.71264D+00    |proj g|=  8.64835D+00
At iterate   45    f= -3.72092D+00    |proj g|=  4.25377D-01
At iterate   50    f= -3.72172D+00    |proj g|=  1.26806D-01

```

* * *

```

Tit   = total number of iterations
Tnf   = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip  = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

* * *

```

      N      Tit      Tnf  Tnint  Skip  Nact      Projg      F
      4       50      91      1      0      0    1.268D-01  -3.722D+00
F = -3.7217246946379170

```

```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT
RUNNING THE L-BFGS-B CODE

```

* * *

```

Machine precision = 2.220D-16
  N =              3      M =              10

```

```

At X0          0 variables are exactly at the bounds

```

```

At iterate    0    f=  2.06552D+03    |proj g|=  7.07257D+05

```

This problem is unconstrained.

```

At iterate    5    f= -8.96283D-01    |proj g|=  4.85152D+00

```

```

At iterate   10    f= -3.02319D+00    |proj g|=  3.76053D+00

```

```

At iterate   15    f= -3.37693D+00    |proj g|=  1.60420D+01

```

```

At iterate   20    f= -3.39220D+00    |proj g|=  3.20094D-01

```

```

At iterate   25    f= -3.39258D+00    |proj g|=  3.31674D-02

```

Warning: more than 10 function and gradient evaluations in the last line search. Termination may possibly be caused by a bad search direction. This problem is unconstrained.

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
3	28	76	1	0	0	7.452D-02	-3.393D+00

F = -3.3926002321547912

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 4 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.06396D+03 |proj g|= 7.06456D+05

At iterate 5 f= -1.54811D+00 |proj g|= 1.15989D+01

At iterate 10 f= -3.52712D+00 |proj g|= 2.35663D+00

At iterate 15 f= -3.70347D+00 |proj g|= 4.59811D+00

At iterate 20 f= -3.71810D+00 |proj g|= 2.86669D-01

At iterate 25 f= -3.71975D+00 |proj g|= 2.99539D-01

At iterate 30 f= -3.71976D+00 |proj g|= 1.79169D-01
 ys=-3.197E-04 -gs= 3.565E-08 BFGS update SKIPPED

Bad direction in the line search;
 refresh the lbfgs memory and restart the iteration.

Warning: more than 10 function and gradient evaluations in the last line search. Termination may possibly be caused by a bad search direction. This problem is unconstrained.

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
4	31	135	2	1	0	1.792D-01	-3.720D+00

F = -3.7197590437067625

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 5 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.16971D+03 |proj g|= 7.41441D+05
 ys=-3.430E-01 -gs= 1.111E+00 BFGS update SKIPPED

At iterate 5 f= -4.93525D-01 |proj g|= 2.88149D+00
 ys=-5.171E+00 -gs= 9.919E-01 BFGS update SKIPPED

At iterate 10 f= -3.25051D+00 |proj g|= 3.12147D+01

At iterate 15 f= -3.30131D+00 |proj g|= 2.41104D+00

At iterate 20 f= -3.69097D+00 |proj g|= 2.01094D+01

At iterate 25 f= -3.70209D+00 |proj g|= 2.94193D+00

At iterate 30 f= -3.72044D+00 |proj g|= 1.22724D+00

At iterate 35 f= -3.72142D+00 |proj g|= 8.69372D-03

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
5	35	94	1	2	0	8.694D-03	-3.721D+00

F = -3.7214177954169956

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH

Summary of best ARMA model for ABC:

SARIMAX Results

```
=====
Dep. Variable:          close    No. Observations:          1085
Model:                 SARIMAX(1, 0, 0)    Log Likelihood          4037.479
Date:                 Sun, 26 Mar 2023    AIC                    -8070.958
Time:                 15:52:55           BIC                    -8060.980
Sample:               0              HQIC                    -8067.181
                        - 1085
```

Covariance Type: opg

```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.9998         0.000    3169.397      0.000         0.999         1.000
sigma2        3.399e-05     4.89e-07     69.451      0.000        3.3e-05        3.5e-05
=====
```

```
=====
Ljung-Box (L1) (Q):                0.10    Jarque-Bera (JB):                114
74.03
Prob(Q):                          0.75    Prob(JB):
0.00
Heteroskedasticity (H):            5.79    Skew:
-0.35
Prob(H) (two-sided):              0.00    Kurtosis:
18.92
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

MSE score for ABC: 7647.104895228071

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 1 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.03696D+00 |proj g|= 2.14672D-05

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

```

N      Tit      Tnf  Tnint  Skip  Nact      Projg      F
1       1       11      1      0      0    2.145D-05  1.037D+00
F =    1.0369568733519654
```

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH

Warning: more than 10 function and gradient
evaluations in the last line search. Termination
may possibly be caused by a bad search direction.
This problem is unconstrained.
This problem is unconstrained.
This problem is unconstrained.

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.03696D+00 |proj g|= 9.96913D-01

At iterate 5 f= 3.84556D-01 |proj g|= 2.75650D-02

At iterate 10 f= 3.59754D-01 |proj g|= 4.65689D-03

At iterate 15 f= 3.57967D-01 |proj g|= 1.80161D-03

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
2	17	24	1	0	0	3.366D-07	3.580D-01

F = 0.35796719152715345

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 3 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 5.99922D-01 |proj g|= 1.29192D+00

At iterate 5 f= -1.74274D-01 |proj g|= 2.06233D+00

At iterate 10 f= -2.31865D-01 |proj g|= 6.77435D-02

At iterate 15 f= -2.32848D-01 |proj g|= 3.17593D-03

```
Bad direction in the line search;  
  refresh the lbfgs memory and restart the iteration.  
  
Line search cannot locate an adequate point after MAXLS  
  function and gradient evaluations.  
  Previous x, f and g restored.  
Possible causes: 1 error in function or gradient evaluation;  
                  2 rounding error dominate computation.  
This problem is unconstrained.
```

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
3	18	64	2	0	0	1.378D-04	-2.329D-01

F = -0.23285267762358666

ABNORMAL_TERMINATION_IN_LNSRCH
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.43470D+03 |proj g|= 2.25674D+05

At iterate 5 f= -1.28235D+00 |proj g|= 2.89243D+00

At iterate 10 f= -2.60585D+00 |proj g|= 1.60212D+00

At iterate 15 f= -2.89850D+00 |proj g|= 1.65832D+00

At iterate 20 f= -2.93842D+00 |proj g|= 3.22398D-01

At iterate 25 f= -2.94163D+00 |proj g|= 2.93106D-01

At iterate 30 f= -2.94177D+00 |proj g|= 6.20897D-04

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
2	32	62	1	0	0	1.359D-06	-2.942D+00

F = -2.9417651716849273

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 3 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.38190D+03 |proj g|= 2.17125D+05

This problem is unconstrained.

At iterate 5 f= -1.58573D+00 |proj g|= 9.08294D+00

At iterate 10 f= -2.34375D+00 |proj g|= 4.34504D+00

At iterate 15 f= -2.74455D+00 |proj g|= 3.19169D+01

At iterate 20 f= -2.92021D+00 |proj g|= 3.99404D+00

At iterate 25 f= -2.94025D+00 |proj g|= 7.25078D-01

At iterate 30 f= -2.94184D+00 |proj g|= 1.93504D-01

At iterate 35 f= -2.94193D+00 |proj g|= 1.10083D-02

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
3	36	63	1	0	0	8.967D-03	-2.942D+00

F = -2.9419272077431575

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 4 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.58913D+03 |proj g|= 2.49815D+05

This problem is unconstrained.


```

At iterate    5    f= -9.45637D-01    |proj g|=  7.85781D+00
At iterate   10    f= -1.64220D+00    |proj g|=  6.52824D-01
At iterate   15    f= -2.43758D+00    |proj g|=  2.43580D+01
At iterate   20    f= -2.53114D+00    |proj g|=  8.18867D-01
At iterate   25    f= -2.56314D+00    |proj g|=  5.38525D+00
At iterate   30    f= -2.93188D+00    |proj g|=  1.05862D+01
At iterate   35    f= -2.93856D+00    |proj g|=  1.18276D+00
At iterate   40    f= -2.94056D+00    |proj g|=  1.93773D+00
At iterate   45    f= -2.94125D+00    |proj g|=  2.29668D+00
At iterate   50    f= -2.94412D+00    |proj g|=  3.20240D-01

```

* * *

```

Tit   = total number of iterations
Tnf   = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip  = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

* * *

```

      N      Tit      Tnf  Tnint  Skip  Nact      Projg      F
      4       50       77      1      0      0    3.202D-01  -2.944D+00
F = -2.9441204535127010

```

```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT
RUNNING THE L-BFGS-B CODE

```

* * *

```
Machine precision = 2.220D-16
```

```
N =          3      M =          10
```

```
At X0          0 variables are exactly at the bounds
```

```
At iterate    0    f=  1.43383D+03    |proj g|=  2.25468D+05
```

```
At iterate    5    f= -1.82243D+00    |proj g|=  4.85574D+00
```

```
This problem is unconstrained.
```

```
At iterate   10    f= -2.59310D+00    |proj g|=  1.14735D+00
At iterate   15    f= -2.60434D+00    |proj g|=  8.23307D-03
At iterate   20    f= -2.60449D+00    |proj g|=  4.45710D-01
At iterate   25    f= -2.88607D+00    |proj g|=  2.33475D+01
At iterate   30    f= -2.93752D+00    |proj g|=  7.99435D-01
At iterate   35    f= -2.94177D+00    |proj g|=  1.27566D-02
```

Warning: more than 10 function and gradient
evaluations in the last line search. Termination
may possibly be caused by a bad search direction.
This problem is unconstrained.

At iterate 40 f= -2.94190D+00 |proj g|= 6.75818D-03

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Proyg	F
3	40	95	1	0	0	6.758D-03	-2.942D+00

F = -2.9419032219365810

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16
 N = 4 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.43868D+03 |proj g|= 2.26610D+05

At iterate 5 f= -1.36299D+00 |proj g|= 2.16641D+01

At iterate 10 f= -2.81795D+00 |proj g|= 3.12833D+01

At iterate 15 f= -2.91403D+00 |proj g|= 1.56683D+00

At iterate 20 f= -2.93353D+00 |proj g|= 6.60428D-01

At iterate 25 f= -2.93732D+00 |proj g|= 6.06072D-02

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Proyg	F
4	27	61	1	0	0	3.113D-02	-2.937D+00

F = -2.9374348334511433

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
 RUNNING THE L-BFGS-B CODE

* * *

```
Machine precision = 2.220D-16
```

```
N =          5      M =          10
```

```
At X0          0 variables are exactly at the bounds
```

```
At iterate    0      f=  1.70866D+03      |proj g|=  2.68572D+05
```

```
This problem is unconstrained.
```

ys=-1.606E+00 -gs= 9.784E-01 BFGS update SKIPPED

At iterate 5 f= -2.35801D+00 |proj g|= 5.37861D-01
 At iterate 10 f= -2.76840D+00 |proj g|= 4.16295D+00
 At iterate 15 f= -2.92387D+00 |proj g|= 1.44228D-01
 At iterate 20 f= -2.92394D+00 |proj g|= 1.02472D+00
 At iterate 25 f= -2.92881D+00 |proj g|= 7.03905D+00
 At iterate 30 f= -2.93725D+00 |proj g|= 1.43219D+00

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Proyg	F
5	31	57	1	1	0	1.432D+00	-2.937D+00

F = -2.9372527918357996

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH

Summary of best ARMA model for AAP:

SARIMAX Results

Dep. Variable:	close	No. Observations:	1085
Model:	SARIMAX(1, 0, 2)	Log Likelihood	3194.371
Date:	Sun, 26 Mar 2023	AIC	-6380.741
Time:	15:52:59	BIC	-6360.784
Sample:	0	HQIC	-6373.186
	- 1085		

Covariance Type: opg

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.9997	0.000	2120.353	0.000	0.999	1.001
ma.L1	0.0133	0.033	0.404	0.687	-0.051	0.078
ma.L2	-0.0635	0.027	-2.345	0.019	-0.117	-0.010
sigma2	0.0002	1.63e-06	99.241	0.000	0.000	0.000

Ljung-Box (L1) (Q):	0.01	Jarque-Bera (JB):	601
79.77			
Prob(Q):	0.92	Prob(JB):	
0.00			
Heteroskedasticity (H):	2.21	Skew:	
-0.30			
Prob(H) (two-sided):	0.00	Kurtosis:	
39.48			

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

MSE score for AAP: 10768.315414664408

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 1 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= -1.91402D-02 |proj g|= 1.77447D-04

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
1	1	15	1	0	0	1.774D-04	-1.914D-02

F = -1.9140238166426338E-002

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= -3.87943D-01 |proj g|= 3.12543D+00

At iterate 5 f= -6.95944D-01 |proj g|= 1.93678D-01

This problem is unconstrained.

Warning: more than 10 function and gradient evaluations in the last line search. Termination may possibly be caused by a bad search direction. This problem is unconstrained.

At iterate 10 f= -6.98657D-01 |proj g|= 1.41320D-02

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Provg	F
2	14	26	1	0	0	3.607D-04	-6.987D-01

F = -0.69866026150373650

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16
 N = 3 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= -4.54065D-01 |proj g|= 3.69505D+00

At iterate 5 f= -1.21580D+00 |proj g|= 5.30011D+00

This problem is unconstrained.

```

At iterate   10    f= -1.28521D+00    |proj g|=  1.56756D-01
At iterate   15    f= -1.28527D+00    |proj g|=  5.63986D-02
At iterate   20    f= -1.28589D+00    |proj g|=  6.63548D-01
At iterate   25    f= -1.28768D+00    |proj g|=  9.61065D-02
At iterate   30    f= -1.29070D+00    |proj g|=  2.04264D+00
At iterate   35    f= -1.30609D+00    |proj g|=  7.09093D-02

```

* * *

```

Tit   = total number of iterations
Tnf   = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip  = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

* * *

```

      N      Tit      Tnf  Tnint  Skip  Nact      Projg      F
      3       39       54       1       0       0  1.208D-06 -1.306D+00
F = -1.3061005256212552

```

```

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL
RUNNING THE L-BFGS-B CODE

```

* * *

```

Machine precision = 2.220D-16

```

```

N =          2      M =          10

```

```

At X0          0 variables are exactly at the bounds

```

```

At iterate    0    f=  1.12697D+03    |proj g|=  4.51851D+05

```

```

At iterate    5    f= -2.13262D+00    |proj g|=  3.42494D+01

```

This problem is unconstrained.


```

At iterate   10    f= -3.42147D+00    |proj g|=  4.02662D+01
At iterate   15    f= -3.85206D+00    |proj g|=  1.46903D+01
At iterate   20    f= -3.87572D+00    |proj g|=  5.21264D-01
At iterate   25    f= -3.87764D+00    |proj g|=  7.09505D-02
At iterate   30    f= -3.87776D+00    |proj g|=  1.33695D-04

```

* * *

```

Tit  = total number of iterations
Tnf  = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

* * *

```

      N      Tit      Tnf  Tnint  Skip  Nact      Projg      F
      2       31       64      1      0      0   9.041D-05  -3.878D+00
F = -3.8777611827802900

```

```

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
RUNNING THE L-BFGS-B CODE

```

* * *

```

Machine precision = 2.220D-16
N =          3      M =          10

```

```

At X0          0 variables are exactly at the bounds

```

```

At iterate    0    f=  1.18037D+03    |proj g|=  4.72760D+05

```

This problem is unconstrained.

```

ys=-6.762E-02  -gs= 7.131E-01 BFGS update SKIPPED

At iterate    5    f= -3.10550D-01    |proj g|=  6.48206D+00
At iterate   10    f= -3.13491D+00    |proj g|=  3.31099D+01
At iterate   15    f= -3.38530D+00    |proj g|=  1.36927D+01
At iterate   20    f= -3.77008D+00    |proj g|=  2.20387D+01
At iterate   25    f= -3.86901D+00    |proj g|=  3.06482D+00
At iterate   30    f= -3.87710D+00    |proj g|=  1.40526D+00
At iterate   35    f= -3.87800D+00    |proj g|=  6.36884D-02
At iterate   40    f= -3.87802D+00    |proj g|=  6.09147D-02

* * *

Tit   = total number of iterations
Tnf   = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip  = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

```

* * *

      N      Tit      Tnf  Tnint  Skip  Nact      Projg      F
      3       40       75      1      1      0    6.091D-02  -3.878D+00
F = -3.8780210437723523

```

```

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
RUNNING THE L-BFGS-B CODE

```

```

* * *

Machine precision = 2.220D-16
  N =           4      M =           10

At X0           0 variables are exactly at the bounds

At iterate    0    f=  1.24945D+03    |proj g|=  5.00477D+05
This problem is unconstrained.

```

```

At iterate    5    f= -1.04323D+00    |proj g|=  1.33284D+00
At iterate   10    f= -2.73702D+00    |proj g|=  1.72254D+01
At iterate   15    f= -2.78273D+00    |proj g|=  2.95495D+00
At iterate   20    f= -3.33571D+00    |proj g|=  1.40761D+01
At iterate   25    f= -3.82294D+00    |proj g|=  1.04256D+01
At iterate   30    f= -3.84325D+00    |proj g|=  6.91221D-01
At iterate   35    f= -3.84846D+00    |proj g|=  3.40007D+00
At iterate   40    f= -3.87825D+00    |proj g|=  4.20006D-01
At iterate   45    f= -3.87831D+00    |proj g|=  8.28653D-01
At iterate   50    f= -3.87839D+00    |proj g|=  1.13850D-02

```

* * *

```

Tit   = total number of iterations
Tnf   = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip  = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

* * *

```

      N      Tit      Tnf  Tnint  Skip  Nact      Projg      F
      4       50      95      1      0      0    1.138D-02  -3.878D+00
F = -3.8783850840965286

```

```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT
RUNNING THE L-BFGS-B CODE

```

* * *

```

Machine precision = 2.220D-16

```

```

N =          3      M =          10

```

```

At X0          0 variables are exactly at the bounds

```

```

At iterate    0    f=  1.12672D+03    |proj g|=  4.51698D+05
ys=-7.228E+00  -gs= 1.834E+00 BFGS update SKIPPED

```

This problem is unconstrained.

```
At iterate    5    f= -2.48001D+00    |proj g|=  9.07147D+00
At iterate   10    f= -3.49994D+00    |proj g|=  5.91187D+00
At iterate   15    f= -3.55941D+00    |proj g|=  4.86307D-01
At iterate   20    f= -3.56000D+00    |proj g|=  2.28403D-02
At iterate   25    f= -3.56012D+00    |proj g|=  2.87268D-01
At iterate   30    f= -3.84374D+00    |proj g|=  4.95196D+01
At iterate   35    f= -3.87245D+00    |proj g|=  2.27184D+00
At iterate   40    f= -3.87749D+00    |proj g|=  9.56717D-01
At iterate   45    f= -3.87800D+00    |proj g|=  3.29201D-02
```

Warning: more than 10 function and gradient
evaluations in the last line search. Termination
may possibly be caused by a bad search direction.
This problem is unconstrained.

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
3	47	103	1	1	0	2.769D-01	-3.878D+00

F = -3.8780073923270884

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 4 M = 10

At X0 0 variables are exactly at the bounds

At iterate	0	f=	1.12587D+03	proj g =	4.51193D+05
At iterate	5	f=	-1.65889D+00	proj g =	3.02171D+00
At iterate	10	f=	-3.12174D+00	proj g =	2.84949D+01
At iterate	15	f=	-3.71812D+00	proj g =	2.82372D+01
At iterate	20	f=	-3.86495D+00	proj g =	1.84096D+00
At iterate	25	f=	-3.87527D+00	proj g =	2.00387D+00
At iterate	30	f=	-3.87831D+00	proj g =	7.39882D-01
At iterate	35	f=	-3.87838D+00	proj g =	1.27649D-01

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
4	36	107	1	0	0	1.250D-01	-3.878D+00

F = -3.8783841963015897

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH
 RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 5 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.12855D+03 |proj g|= 4.52799D+05

Warning: more than 10 function and gradient
evaluations in the last line search. Termination
may possibly be caused by a bad search direction.
This problem is unconstrained.

```

At iterate    5    f= -1.34035D+00    |proj g|=  1.22484D+01
At iterate   10    f= -3.13153D+00    |proj g|=  3.17954D+01
At iterate   15    f= -3.32140D+00    |proj g|=  4.03264D-01
At iterate   20    f= -3.33449D+00    |proj g|=  1.16754D+01
At iterate   25    f= -3.86968D+00    |proj g|=  3.57769D+00
At iterate   30    f= -3.87016D+00    |proj g|=  4.21245D-01
At iterate   35    f= -3.87393D+00    |proj g|=  3.04728D-01
At iterate   40    f= -3.87431D+00    |proj g|=  3.60578D-01
At iterate   45    f= -3.87431D+00    |proj g|=  6.10923D-02

```

* * *

```

Tit   = total number of iterations
Tnf   = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip  = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

* * *

```

      N      Tit      Tnf  Tnint  Skip  Nact      Projg      F
      5       47     100      1      0      0    5.422D-02  -3.874D+00
F = -3.8743102397723996

```

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH

Summary of best ARMA model for ABBV:

SARIMAX Results

```

=====
Dep. Variable:          close    No. Observations:          1085
Model:                SARIMAX(1, 0, 0)    Log Likelihood          4207.371
Date:                Sun, 26 Mar 2023    AIC                   -8410.742
Time:                15:53:05            BIC                   -8400.763
Sample:                0                HQIC                  -8406.964
                    - 1085

```

Covariance Type: opg

```

=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.9997          0.000    2085.045      0.000          0.999          1.001
sigma2         2.486e-05      6.11e-07     40.665      0.000      2.37e-05      2.61e-05
=====

```

```

=====
Ljung-Box (L1) (Q):                0.60    Jarque-Bera (JB):                8
15.73
Prob(Q):                0.44    Prob(JB):
0.00
Heteroskedasticity (H):            1.50    Skew:
-0.35
Prob(H) (two-sided):            0.00    Kurtosis:
7.19

```

```
=====
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

MSE score for ABBV: 7659.775864189195

Scoring ARMA model - from the output above, there is a best model for each stock such as SARIMAX(1, 0, 0) for ABBV. The MSE for each model is below.

```
In [30]: from sklearn.metrics import mean_squared_error
import warnings
from statsmodels.tools.sm_exceptions import ConvergenceWarning
warnings.simplefilter('ignore', ConvergenceWarning)
warnings.filterwarnings("ignore")
# Create a dictionary to store the MSE scores for each stock
mse_scores = {}

# Loop through each stock in the dataframe
for stock in df_test['name'].unique():

    # Filter the dataframe to only include the current stock
    stock_df = df_test[df_test['name'] == stock]['close']

    # Use the best ARMA model to predict the test data
    arma_model = arma_models.get(stock)
    if arma_model is None:
        print(f"No ARMA model found for {stock}. Skipping.")
        continue
    pred = arma_model.predict(start=len(df_train), end=len(df_train)+len(stock_

    # Calculate the MSE score and store it in the dictionary
    mse_score = mean_squared_error(stock_df, pred)
    mse_scores[stock] = mse_score

    # Print the MSE score
    print(f"MSE score for {stock}: {mse_score}")
```

MSE score for AAL: 2500.3323356590035

MSE score for AAPL: 25897.155535940266

MSE score for AAP: 10768.315414664408

MSE score for ABBV: 7659.775864189195

MSE score for ABC: 7647.104895228071

ARMA:

Summary of best ARMA model for AAL: SARIMAX(2, 0, 0)

Summary of best ARMA model for AAPL: SARIMAX(1, 0, 0)

Summary of best ARMA model for ABC: SARIMAX(1, 0, 0)

Summary of best ARMA model for AAP: SARIMAX(1, 0, 2)

Summary of best ARMA model for ABBV: SARIMAX(1, 0, 0)

MSE score for AAL: 2500.3323356590035

MSE score for AAPL: 25897.155535940266

MSE score for AAP: 10768.315414664408

MSE score for ABBV: 7659.775864189195

MSE score for ABC: 7647.104895228071

```
In [31]: #Model 2 - ARIMA
import statsmodels.api as sm
import itertools
from sklearn.metrics import mean_squared_error

import warnings
from statsmodels.tools.sm_exceptions import ConvergenceWarning
warnings.simplefilter('ignore', ConvergenceWarning)
warnings.filterwarnings("ignore")

# Create a dictionary to store the ARIMA models for each stock
arma_models = {}

# Define the orders to test for the ARIMA model
p = range(0, 3)
d = range(0, 3)
q = range(0, 3)
orders = list(itertools.product(p, d, q))

# Loop through each stock in the dataframe
for stock in df_train['name'].unique():

    # Filter the dataframe to only include the current stock
    stock_df = df_train[df_train['name'] == stock]['close']

    # Train multiple ARIMA models with different orders
    best_aic = float('inf')
    best_order = None
    best_model = None

    for order in orders:
        try:
            arma_model = sm.tsa.ARIMA(stock_df, order=order).fit()
            if arma_model.aic < best_aic:
                best_aic = arma_model.aic
                best_order = order
                best_model = arma_model
        except Exception as e:
            print(f"Error fitting ARIMA model for {stock} with order {order}: {e}")
            continue

    # Store the best model in the dictionary with the stock name as the key
    arma_models[stock] = best_model

# Print the summary of the best model
if best_model is not None:
    print(f"Summary of best ARIMA model for {stock}:")
    print(best_model.summary())
```

```
# Predict the test data using the best model
if stock in df_test['name'].unique():
    test_stock_df = df_test[df_test['name'] == stock]['close']
    pred = best_model.predict(start=len(df_train), end=len(df_train)+len(test_stock_df))
    mse_score = mean_squared_error(test_stock_df, pred)
    print(f"MSE score for {stock}: {mse_score}")
```

Summary of best ARIMA model for AAL:

SARIMAX Results

```

=====
Dep. Variable:          close    No. Observations:          1085
Model:                ARIMA(1, 1, 0)    Log Likelihood          4298.886
Date:                Sun, 26 Mar 2023    AIC                    -8593.771
Time:                15:53:12    BIC                    -8583.794
Sample:                0    HQIC                    -8589.994
                        - 1085
Covariance Type:          opg
=====

```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.0617	0.027	2.316	0.021	0.009	0.114
sigma2	2.103e-05	6.21e-07	33.865	0.000	1.98e-05	2.22e-05

```

=====
Ljung-Box (L1) (Q):          0.01    Jarque-Bera (JB):          2
53.65
Prob(Q):          0.92    Prob(JB):
0.00
Heteroskedasticity (H):      2.04    Skew:
-0.27
Prob(H) (two-sided):      0.00    Kurtosis:
5.31
=====
=====

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

MSE score for AAL: 2482.556037403676

Summary of best ARIMA model for AAPL:

SARIMAX Results

```

=====
Dep. Variable:          close    No. Observations:          1085
Model:                ARIMA(0, 1, 0)    Log Likelihood          3713.184
Date:                Sun, 26 Mar 2023    AIC                    -7424.368
Time:                15:53:20    BIC                    -7419.379
Sample:                0    HQIC                    -7422.479
                        - 1085
Covariance Type:          opg
=====

```

	coef	std err	z	P> z	[0.025	0.975]
sigma2	6.192e-05	1.6e-06	38.799	0.000	5.88e-05	6.5e-05

```

=====
Ljung-Box (L1) (Q):          0.39    Jarque-Bera (JB):          5
84.99
Prob(Q):          0.53    Prob(JB):
0.00
Heteroskedasticity (H):      1.86    Skew:
-0.12
Prob(H) (two-sided):      0.00    Kurtosis:
6.59
=====
=====

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

MSE score for AAPL: 25812.018824137722

Summary of best ARIMA model for ABC:

SARIMAX Results

```
=====
Dep. Variable:          close    No. Observations:          1085
Model:                ARIMA(0, 1, 0)    Log Likelihood          4037.741
Date:                Sun, 26 Mar 2023    AIC                    -8073.481
Time:                15:53:29          BIC                    -8068.493
Sample:                0              HQIC                    -8071.593
                        - 1085
Covariance Type:                opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
sigma2	3.402e-05	4.89e-07	69.524	0.000	3.31e-05	3.5e-05

```
=====
Ljung-Box (L1) (Q):                0.10    Jarque-Bera (JB):                114
43.46
Prob(Q):                0.76    Prob(JB):
0.00
Heteroskedasticity (H):            5.82    Skew:
-0.35
Prob(H) (two-sided):            0.00    Kurtosis:
18.90
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

MSE score for ABC: 7610.837569265233

Summary of best ARIMA model for AAP:

SARIMAX Results

```
=====
Dep. Variable:          close    No. Observations:          1085
Model:                ARIMA(2, 1, 2)    Log Likelihood          3198.191
Date:                Sun, 26 Mar 2023    AIC                    -6386.383
Time:                15:53:42          BIC                    -6361.441
Sample:                0              HQIC                    -6376.940
                        - 1085
Covariance Type:                opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.3462	0.211	-1.639	0.101	-0.760	0.068
ar.L2	-0.6963	0.183	-3.808	0.000	-1.055	-0.338
ma.L1	0.3743	0.224	1.671	0.095	-0.065	0.813
ma.L2	0.6361	0.196	3.250	0.001	0.253	1.020
sigma2	0.0002	1.83e-06	87.522	0.000	0.000	0.000

```
=====
Ljung-Box (L1) (Q):                0.10    Jarque-Bera (JB):                595
12.71
Prob(Q):                0.76    Prob(JB):
0.00
Heteroskedasticity (H):            2.22    Skew:
-0.30
```

Prob(H) (two-sided): 0.00 Kurtosis:
39.29

=====

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

MSE score for AAP: 10664.862358570012

Summary of best ARIMA model for ABBV:

SARIMAX Results

```
=====
Dep. Variable:      close    No. Observations:      1085
Model:              ARIMA(0, 1, 0)    Log Likelihood      4207.211
Date:              Sun, 26 Mar 2023    AIC      -8412.422
Time:              15:53:54    BIC      -8407.433
Sample:            0    HQIC      -8410.533
                   - 1085
Covariance Type:    opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
sigma2	2.489e-05	6.1e-07	40.816	0.000	2.37e-05	2.61e-05

```
=====
Ljung-Box (L1) (Q):      0.58    Jarque-Bera (JB):      8
12.63
Prob(Q):      0.44    Prob(JB):
0.00
Heteroskedasticity (H):      1.49    Skew:
-0.35
Prob(H) (two-sided):      0.00    Kurtosis:
7.18
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

MSE score for ABBV: 7626.947322373217

For example, Summary of best ARIMA model for AAL:ARIMA(1, 1, 0) and MSE score for AAL:
2482.556037403676

ARIMA:

Summary of best ARIMA model for AAL: ARIMA(1, 1, 0)

MSE score for AAL: 2482.556037403676

Summary of best ARIMA model for AAPL: ARIMA(0, 1, 0)

MSE score for AAPL: 25812.018824137722

Summary of best ARIMA model for ABC: ARIMA(0, 1, 0)

MSE score for ABC: 7610.837569265233

Summary of best ARIMA model for AAP: ARIMA(2, 1, 2)

MSE score for AAP: 10664.862358570012

Summary of best ARIMA model for ABBV: ARIMA(0, 1, 0)

MSE score for ABBV: 7626.947322373217

In []:

LSTM

```
In [32]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout
from sklearn.metrics import mean_squared_error

# Define hyperparameters to test
lstm_units = [64, 128]
dropout_rates = [0.1, 0.2, 0.3]
batch_sizes = [32, 64]

# Create a dictionary to store the LSTM models for each stock
lstm_models = {}

# Loop through each stock in the dataframe
for stock in df_train['name'].unique():
    print(f"Training LSTM models for {stock}...")

    # Filter the dataframe to only include the current stock
    stock_df = df_train[df_train['name'] == stock]['close']

    # Prepare data for LSTM model
    x_train, y_train = [], []
    for i in range(60, len(stock_df)):
        x_train.append(stock_df[i-60:i])
        y_train.append(stock_df[i])
    x_train, y_train = np.array(x_train), np.array(y_train)

    # Train multiple LSTM models with different hyperparameters
    best_mse = float('inf')
    best_model = None

    for lstm_unit in lstm_units:
        for dropout_rate in dropout_rates:
            for batch_size in batch_sizes:
                # Define LSTM model
                model = Sequential()
                model.add(LSTM(lstm_unit, input_shape=(60, 1)))
                model.add(Dropout(dropout_rate))
                model.add(Dense(1))

                # Compile LSTM model
                model.compile(loss='mean_squared_error', optimizer='adam')

                # Train LSTM model
                model.fit(x_train, y_train, epochs=10, batch_size=batch_size, v

    # Predict the test data using the LSTM model
    if stock in df_test['name'].unique():
        test_stock_df = df_test[df_test['name'] == stock]['close']
        x_test = []
        for i in range(60, len(test_stock_df)):
            x_test.append(test_stock_df[i-60:i])
```

```
x_test = np.array(x_test)
pred = model.predict(x_test)
mse_score = mean_squared_error(test_stock_df[60:], pred)
if mse_score < best_mse:
    best_mse = mse_score
    best_model = model

# Store the best model in the dictionary with the stock name as the key
lstm_models[stock] = best_model

# Print the MSE score for the best LSTM model
if best_model is not None:
    print(f"MSE score for {stock}: {best_mse}")
```

Training LSTM models for AAL...

2023-03-26 15:54:01.440795: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: SSE4.1 SSE4.2 AVX AVX2 FMA
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

```
4/4 [=====] - 1s 12ms/step
4/4 [=====] - 1s 18ms/step
4/4 [=====] - 1s 12ms/step
4/4 [=====] - 0s 11ms/step
4/4 [=====] - 0s 12ms/step
4/4 [=====] - 0s 10ms/step
4/4 [=====] - 1s 25ms/step
4/4 [=====] - 1s 24ms/step
4/4 [=====] - 1s 26ms/step
4/4 [=====] - 1s 25ms/step
4/4 [=====] - 1s 25ms/step
4/4 [=====] - 1s 23ms/step
```

MSE score for AAL: 2192.2762817762896

Training LSTM models for AAPL...

```
4/4 [=====] - 1s 12ms/step
4/4 [=====] - 0s 11ms/step
4/4 [=====] - 0s 11ms/step
4/4 [=====] - 1s 13ms/step
4/4 [=====] - 1s 14ms/step
4/4 [=====] - 1s 21ms/step
4/4 [=====] - 1s 32ms/step
4/4 [=====] - 1s 31ms/step
4/4 [=====] - 1s 24ms/step
4/4 [=====] - 1s 24ms/step
4/4 [=====] - 1s 23ms/step
4/4 [=====] - 1s 24ms/step
```

MSE score for AAPL: 26832.771402801838

Training LSTM models for ABC...

```
4/4 [=====] - 1s 14ms/step
4/4 [=====] - 0s 11ms/step
4/4 [=====] - 0s 10ms/step
4/4 [=====] - 1s 11ms/step
4/4 [=====] - 0s 12ms/step
4/4 [=====] - 0s 10ms/step
4/4 [=====] - 1s 25ms/step
4/4 [=====] - 1s 25ms/step
4/4 [=====] - 1s 24ms/step
4/4 [=====] - 1s 22ms/step
4/4 [=====] - 3s 97ms/step
4/4 [=====] - 1s 26ms/step
```

MSE score for ABC: 7134.804472304175

Training LSTM models for AAP...

```
4/4 [=====] - 1s 23ms/step
4/4 [=====] - 1s 13ms/step
4/4 [=====] - 1s 11ms/step
4/4 [=====] - 0s 10ms/step
4/4 [=====] - 1s 18ms/step
4/4 [=====] - 0s 12ms/step
4/4 [=====] - 1s 31ms/step
4/4 [=====] - 1s 26ms/step
4/4 [=====] - 1s 26ms/step
4/4 [=====] - 1s 26ms/step
4/4 [=====] - 1s 35ms/step
4/4 [=====] - 1s 25ms/step
```

MSE score for AAP: 9292.259072552697

Training LSTM models for ABBV...

```
4/4 [=====] - 1s 15ms/step
4/4 [=====] - 0s 12ms/step
4/4 [=====] - 1s 11ms/step
4/4 [=====] - 0s 11ms/step
```



```

4/4 [=====] - 0s 11ms/step
4/4 [=====] - 0s 10ms/step
4/4 [=====] - 1s 19ms/step
4/4 [=====] - 1s 19ms/step
4/4 [=====] - 1s 22ms/step
4/4 [=====] - 2s 73ms/step
4/4 [=====] - 1s 24ms/step
4/4 [=====] - 1s 38ms/step
MSE score for ABBV: 8621.094947492405

```

The average MSE on the 5 best ARMA models for the 5 stocks was 10894.532, the largest mean squared error for all of the models. The ARIMA models average was 10839.438, the middle value for the three. The LSTM model had an average MSE of 10814.636. This is the lowest average MSE and the best bet for predicting stocks.

The ARMA model is a simple yet powerful time-series model that assumes that the current value of a time series is a linear combination of its past values and random error terms. The ARIMA model is an extension of the ARMA model that includes differencing of the time series to make it stationary. The LSTM model is a type of recurrent neural network that can capture complex and nonlinear patterns in time-series data. LSTMs can learn from the temporal dependencies and long-term memory of a time series by using a memory cell, input, and forget gates.

	best ARMA model	ARMA MSE	best ARIMA model	ARIMA MSE	LSTM MSE
AAL	SARIMAX(2, 0, 0)	2500.33	ARIMA(1, 1, 0)	2482.55	2192.27
AAPL	SARIMAX(1, 0, 0)	25897.15	ARIMA(0, 1, 0)	25812.01	26832.77
AAP	SARIMAX(1, 0, 2)	10768.31	ARIMA(2, 1, 2)	10664.86	9292.25
ABBV	SARIMAX(1, 0, 0)	7659.77	ARIMA(0, 1, 0)	7626.94	8621.09
ABC	SARIMAX(1, 0, 0)	7647.1	ARIMA(0, 1, 0)	7610.83	7134.8
AVG MSE		10894.532		10839.438	10814.636
		biggest MSE		middle MSE	smallest MSE